

Lab #4: Variable Selection and Regularization

CS 109A, STAT 121A, AC 209A: Data Science

Fall 2016

Harvard

Today's lab: Problem 1

- a) Visualize correlation in data
- b) Select minimal subset of predictors
 - Exhaustive search
 - Step-wise forward search


Correlation Matrix

- Pearson correlation coefficient:

- Measure of linear dependence between the predictors i and j

$$\rho_{ij} = \frac{\mathbf{E}[(X_i - \mu_i)(X_j - \mu_j)]}{\sqrt{\mathbf{E}[(X_i - \mu_i)^2]} \sqrt{\mathbf{E}[(X_j - \mu_j)^2]}}$$

Covariance(X_i, X_j)



where $\mu_i = \mathbf{E}[X_i]$, $\mu_j = \mathbf{E}[X_j]$


Correlation Matrix

- Pearson correlation coefficient:

- Measure of linear dependence between the predictors i and j

$$\rho_{ij} = \frac{\mathbf{E}[(X_i - \mu_i)(X_j - \mu_j)]}{\sqrt{\mathbf{E}[(X_i - \mu_i)^2]} \sqrt{\mathbf{E}[(X_j - \mu_j)^2]}}$$

Covariance(X_i, X_j)



where $\mu_i = \mathbf{E}[X_i]$, $\mu_j = \mathbf{E}[X_j]$

- What is the dimension of the correlation matrix?

Solve Part 1(a)

- Compute correlation matrix
 - `np.corrcoef(...)`
- Visualize correlation matrix:
 - `ax.pcolor(...)`

Variable Selection

Exhaustive Search

- For each size 'k':
 - Enumerate all subsets of size 'k'
 - Fit regression model for each subset
 - Pick subset with maximum R^2
- Use BIC to choose best size, and output optimal subset for that size

Enumerating Subsets

- Enumerate all subsets of predictors $\{0, 1, 2, 3\}$
 - Subsets of size 1: $\{0\}, \{1\}, \{2\}, \{3\}$
 - Subsets of size 2: $\{0, 1\}, \{0, 2\}, \{0, 3\},$
 $\{1, 2\}, \{1, 3\}, \{2, 3\}$
 - Subsets of size 3: $\{0, 1, 2\}, \{0, 1, 3\},$
 $\{0, 2, 3\}, \{1, 2, 3\}$
 - Subsets of size 4: $\{0, 1, 2, 3\}$

Enumerating Subsets

Best R^2 within
each group

- Enumerate all subsets of predictors {0, 1, 2, 3}
 - Subsets of size 1: {0}, {1}, {2}, {3} → Best 1-subset
 - Subsets of size 2: {0, 1}, {0, 2}, {0, 3},
{1, 2}, {1, 3}, {2, 3} → Best 2-subset
 - Subsets of size 3: {0, 1, 2}, {0, 1, 3},
{0, 2, 3}, {1, 2, 3} → Best 3-subset
 - Subsets of size 4: {0, 1, 2, 3} → Best 4-subset

Enumerating Subsets

- Enumerate all subsets of predictors $\{0, 1, 2, 3\}$

- Subsets of size 1: $\{0\}, \{1\}, \{2\}, \{3\}$

→ Best 1-subset

- Subsets of size 2: $\{0, 1\}, \{0, 2\}, \{0, 3\},$
 $\{1, 2\}, \{1, 3\}, \{2, 3\}$

→ Best 2-subset

- Subsets of size 3: $\{0, 1, 2\}, \{0, 1, 3\},$
 $\{0, 2, 3\}, \{1, 2, 3\}$

→ Best 3-subset

- Subsets of size 4: $\{0, 1, 2, 3\}$

→ Best 4-subset

***Choose subset
with lowest BIC***

Enumerating Subsets

- Generate all subsets of `set` of size `k`

```
subsets_k = itertools.combinations(set, k)
```

- Output is a list-like object
- Iterating through the generated subsets

```
for subset in subsets_k:
```

```
    ...
```

Putting it together

Outer loop: iterate over sizes 1 d

for k **in** range(d):

Enumerate subsets of size 'k'

subsets_k = itertools.combinations(predictors, k)

Putting it together

Outer loop: iterate over sizes 1 d

for k **in** range(d):

Enumerate subsets of size 'k'

subsets_k = itertools.combinations(predictors, k)

Inner loop: iterate through subsets_k

for subset **in** subsets_k :

Fit regression model using 'subset' and calculate R^2

Keep track of subset with highest R^2

...

Putting it together

Outer loop: iterate over sizes 1 d

for k **in** range(d):

Enumerate subsets of size 'k'

subsets_k = itertools.combinations(predictors, k)

Inner loop: iterate through subsets_k

for subset **in** subsets_k :

Fit regression model using 'subset' and calculate R^2

Keep track of subset with highest R^2

...

Finds
k-sized subset
with best R^2

Putting it together

Outer loop: iterate over sizes 1 d

for k **in** range(d):

Enumerate subsets of size 'k'

subsets_k = itertools.combinations(predictors, k)

Inner loop: iterate through subsets_k

for subset **in** subsets_k :

Fit regression model using 'subset' and calculate R^2

Keep track of subset with highest R^2

...

Finds
k-sized subset
with best R^2

Compute BIC of the subset you get from the inner loop

Compare with lowest BIC so far

Solve Part 1(b)

- Implement exhaustive search
- Implement step-wise forward selection

Step-wise Forward Selection

- Start with empty set
- Repeat for every subset size 1, ..., d:
 - For each predictor **not chosen so far**: add the predictor and fit a regression model
 - **Find predictor that improves the R^2 the most**
- Use **BIC** to choose best subset size

Current & Remaining Lists

- **Outer loop:** iterate over $1, \dots, d$

`current_predictors`

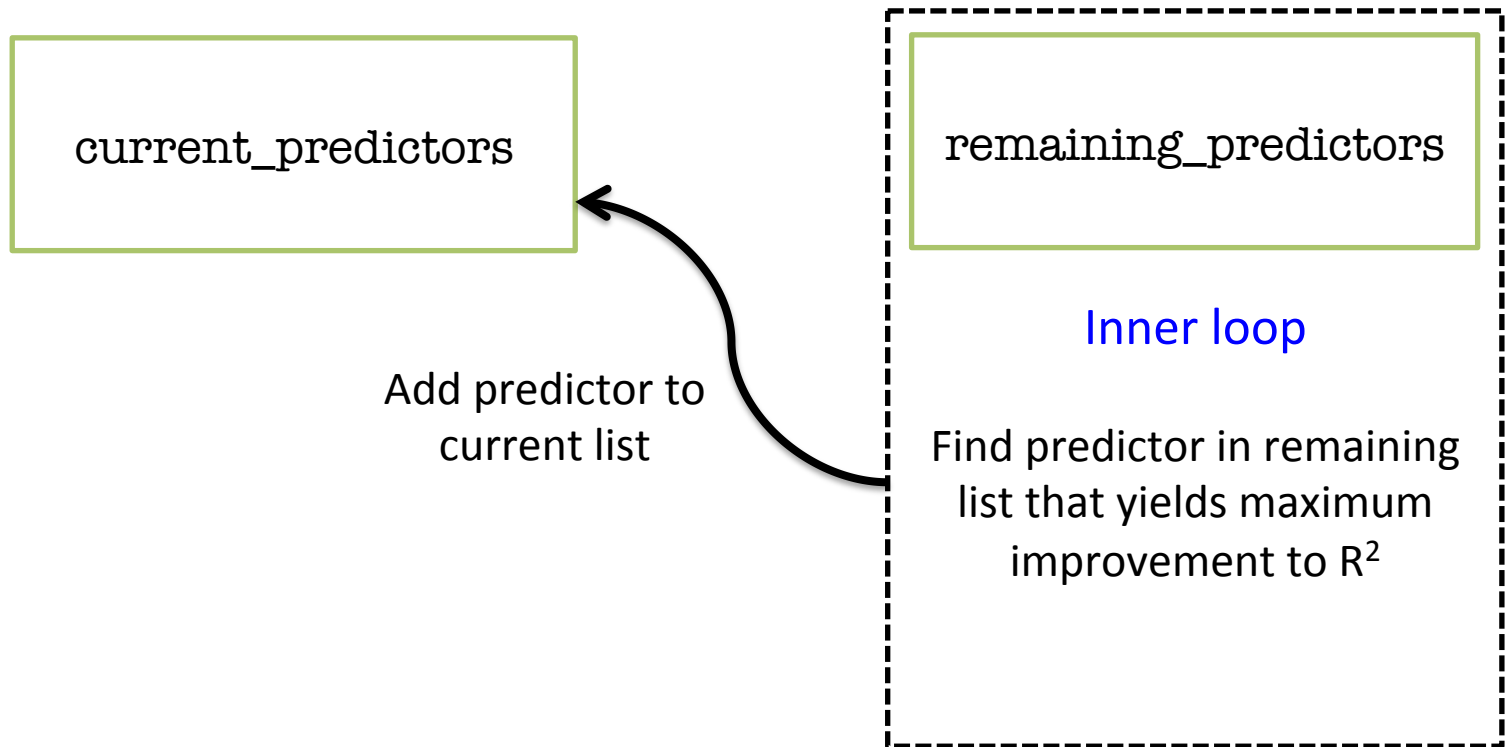
(initially: empty)

`remaining_predictors`

(initially: all predictors)

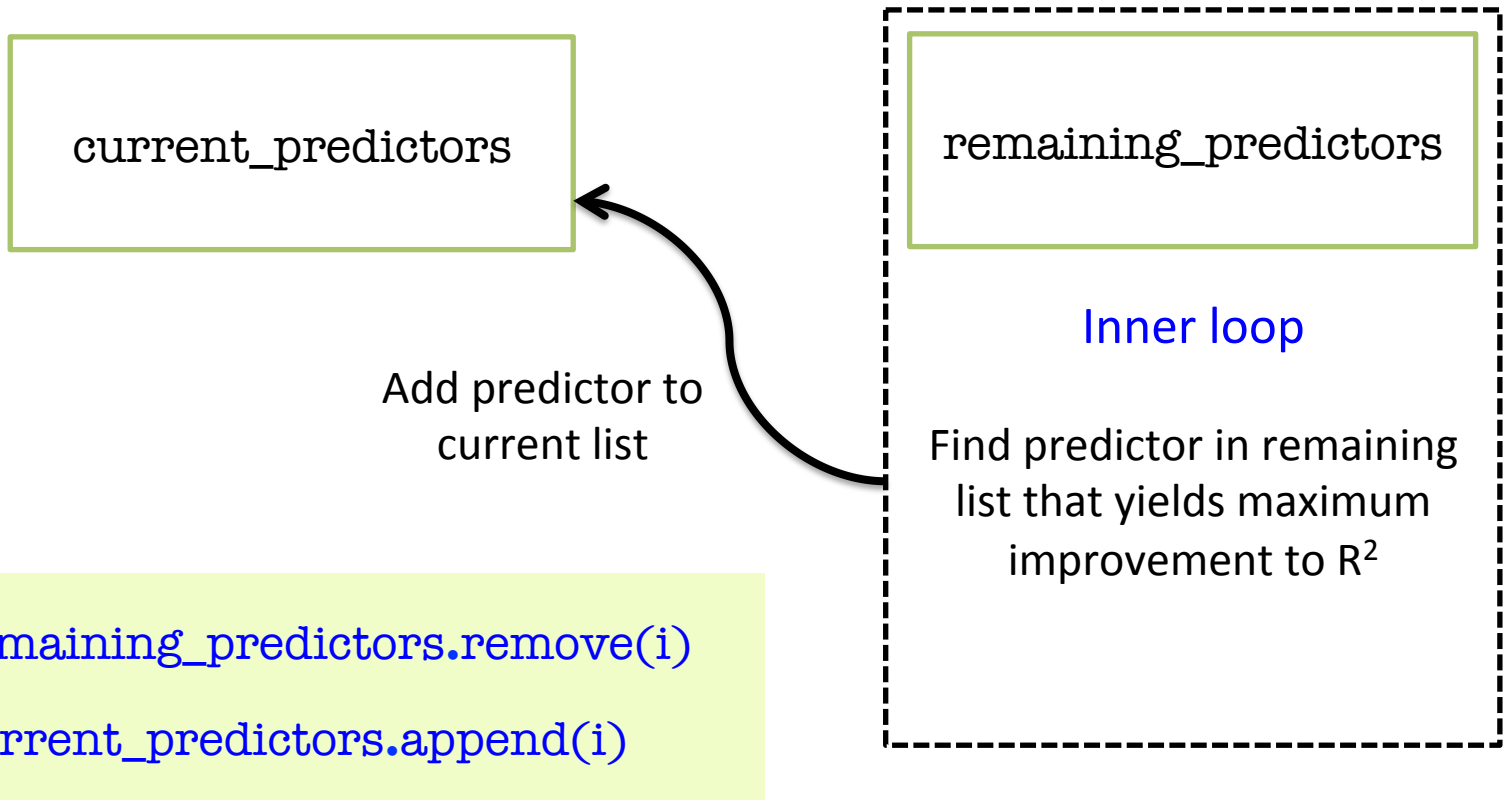
Current & Remaining Lists

- **Outer loop:** iterate over $1, \dots, d$



Current & Remaining Lists

- **Outer loop:** iterate over $1, \dots, d$



Putting it together

```
current_predictors = []  
remaining_predictors = range(d)  
  
# Outer loop: iterate over sizes 1 .... d  
for size in range(d):
```

Putting it together

```
current_predictors = []
```

```
remaining_predictors = range(d)
```

```
# Outer loop: iterate over sizes 1 .... d
```

```
for size in range(d):
```

```
# Inner loop: iterate over remaining_predictors
```

```
for i in remaining_predictors :
```

```
# Make a copy of current_predictors, add 'i' to the copied list
```

```
# Fit regression model to the copied list, evaluate R^2
```

Putting it together

```
current_predictors = []  
remaining_predictors = range(d)
```

Outer loop: iterate over sizes 1 d

```
for size in range(d):
```

Finds predictor
that maximizes
 R^2 the most

Inner loop: iterate over remaining_predictors

```
for i in remaining_predictors :
```

Make a copy of current_predictors, add 'i' to the copied list

Fit regression model to the copied list, evaluate R^2

Putting it together

```
current_predictors = []  
remaining_predictors = range(d)
```

Outer loop: iterate over sizes 1 d

```
for size in range(d):
```

Finds predictor
that maximizes
 R^2 the most

```
# Inner loop: iterate over remaining_predictors
```

```
for i in remaining_predictors :
```

```
# Make a copy of current_predictors, add 'i' to the copied list
```

```
# Fit regression model to the copied list, evaluate  $R^2$ 
```

```
# Add predictor you get to current_predictors
```

```
# Remove the predictor from remaining_predictors
```

```
# Compute BIC of current_predictors, and compare with best BIC so far
```


Solve Part 1(b)

- Implement exhaustive search
- Implement step-wise forward selection

Dealing with Categorical Predictors

- **One-hot encoding:** Binary encoding of categorical predictors
- If predictor Z takes K categories $\{c_1, \dots, c_K\}$, replace it with **K binary predictors** Z_1, \dots, Z_K :
 - $Z_i = 1$ when Z takes value c_i and is 0 otherwise

One-hot Encoding in pandas

- How do you identify categorical attributes?
 - Look for **data type** (string or object)
 - Look for **number of unique values** (< 8?)

One-hot Encoding in pandas

- How do you identify categorical attributes?
 - Look for `data type` (string or object)
 - Look for `number of unique values` (< 8?)
- Transforming single predictor:
 - `pandas.get_dummies(predictor)`
 - Input is a single column as a df
 - Output is a df of multiple binary predictors

One-hot Encoding in pandas

- How do you identify categorical attributes?
 - Look for data type (string or object)
 - Look for number of unique values (< 8?)
- Transforming single predictor:
 - `pandas.get_dummies(predictor)`
 - Input is a single column as a df
 - Output is a df of multiple binary predictors
- Append new predictors to data frame:
`pandas.concat(...)`

Putting it together

BP	Blood type	Height
	O	
	B	
	AB	
	A	
	...	

Putting it together

BP	Blood type	Height
	O	
	B	
	AB	
	A	
	...	

O	A	B	AB
1	0	0	0
0	0	1	0
0	0	0	1
0	1	0	0
...

Putting it together

BP	Blood type	Height
	O	
	B	
	AB	
	A	
	...	

O	A	B	AB
1	0	0	0
0	0	1	0
0	0	0	1
0	1	0	0
...

Start with
empty df,
append one
by one

BP

Putting it together

BP	Blood type	Height
	O	
	B	
	AB	
	A	
	...	

O	A	B	AB
1	0	0	0
0	0	1	0
0	0	0	1
0	1	0	0
...

Start with
empty df,
append one
by one

BP	Blood type			
	O	A	B	AB
	1	0	0	0
	0	0	1	0
	0	0	0	1
	0	1	0	0

Putting it together

BP	Blood type	Height
	O	
	B	
	AB	
	A	
	...	

O	A	B	AB
1	0	0	0
0	0	1	0
0	0	0	1
0	1	0	0
...

Start with
empty df,
append one
by one

BP	Blood type				Height
	O	A	B	AB	
	1	0	0	0	
	0	0	1	0	
	0	0	0	1	
	0	1	0	0	
	

Putting it together

```
# Record start index of attribute in expanded feature vector
start_index = np.zeros(d + 1) # last entry would contain the len of vector +1

# Create a new data frame to store one-hot encoding of attributes
x_df_expanded = pd.DataFrame({})

# Iterate over all attributes
for column in x_df.columns:
    # Check if attribute is categorical: has less than 8 unique values,
    # or is string values (column has type 'object')
    if len(x_df[column].unique()) < 8 or x_df[column].dtype == np.dtype('object'):
        # use one-hot encoding for this column
        encoding = pd.get_dummies(x_df[column])
        # append expanded attribute to data frame
        x_df_expanded = pd.concat([x_df_expanded, encoding], axis=1)
    else:
        x_df_expanded = pd.concat([x_df_expanded, x_df[[column]]], axis=1)
```